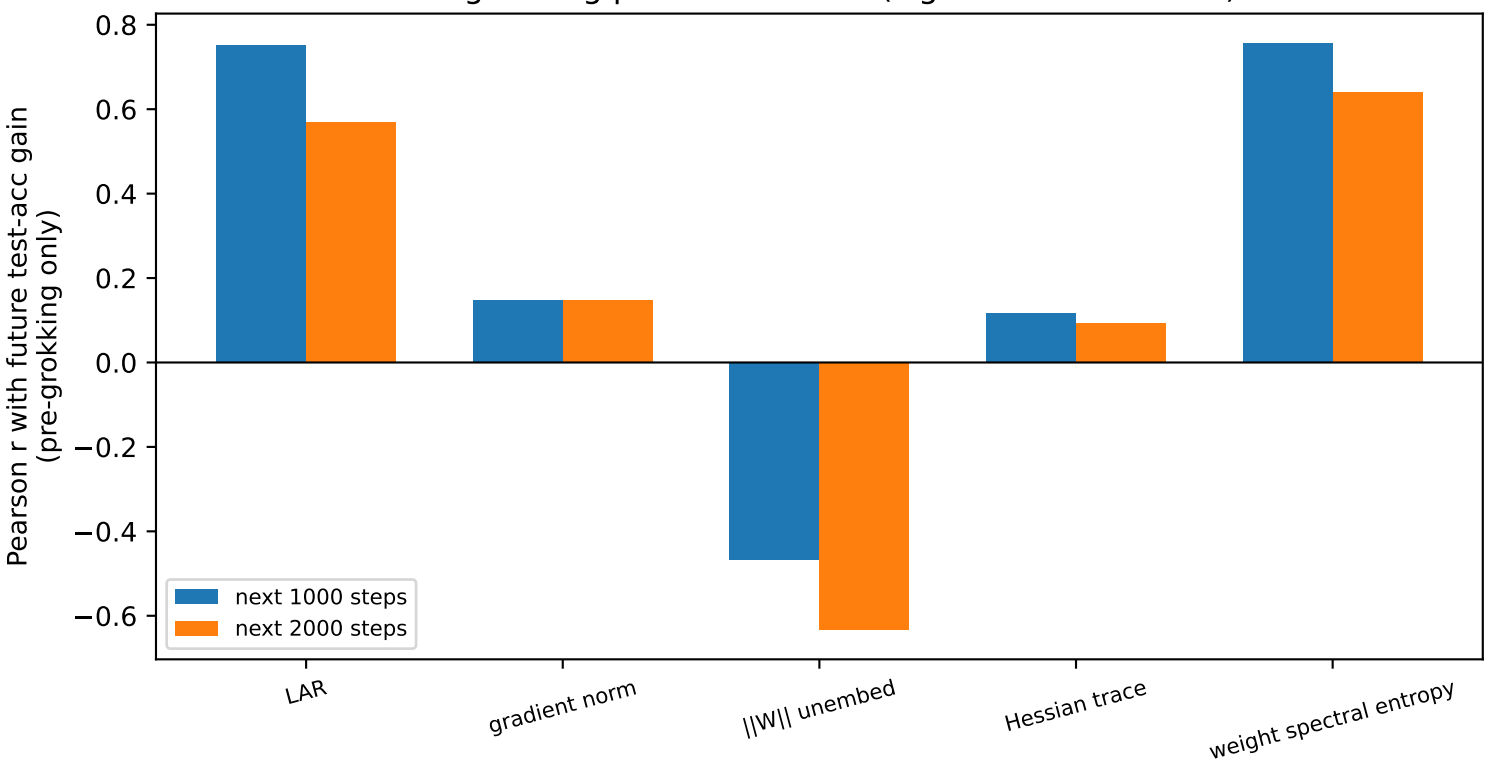
Pre-grokking predictive value (signals use no labels)
Pearson r with future test-acc gain
(pre-grokking only)
0.8
0.6
0.4
0.2
0.0
−0.2
−0.4
−0.6
next 1000 steps
next 2000 steps
LAR
gradient norm
||W|| unembed
Hessian trace
weight spectral entropy