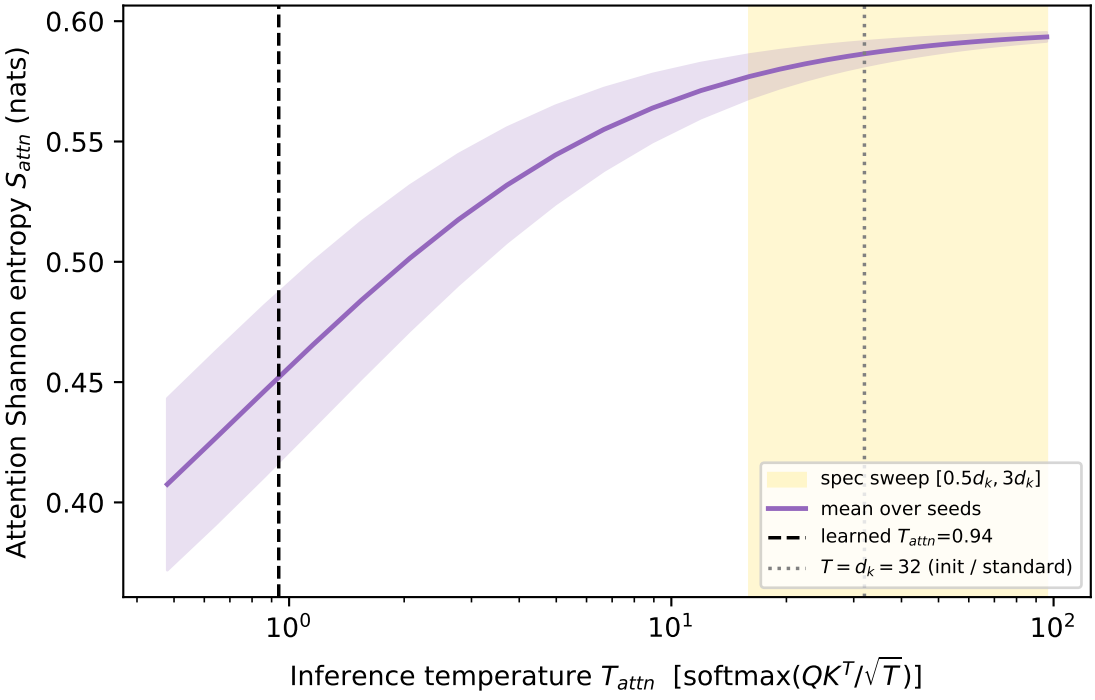
Attention entropy vs inference temperature
Attention Shannon entropy S_{attn} (nats)
0.60
0.55
0.50
0.45
0.40
10^0
10^1
10^2
Inference temperature T_{attn} [softmax($QK^T/\sqrt{T}$)]
spec sweep [$0.5d_k$, $3d_k$]
mean over seeds
learned T_{attn}=0.94
$T = d_k = 32$ (init / standard)